

Big Data systems

From the streets of silicon valley to the 2013 World Tour

Carlo Curino

Agenda

History

2003 cluster as an embedded device

2013 the advent of general purpose computing

Cluster OS (new abstraction)

Mesos / Omega / YARN

Supports for new applications

Apps, services and Meta-frameworks

History

“From the dive bars of silicon valley to the World Tour”

The origin

Purpose-built technology

Within large web companies

Well targeted mission (process webcrawl)

→ scale and fault tolerance



Google leading the pack

Google File System + MapReduce (2003/2004)

Open-source and parallel efforts

Yahoo! Hadoop ecosystem HDFS + MR (2006/2007)

Microsoft Scope/Cosmos (2008) (more than MR)

In-house growth

What was the key to success for Hadoop?

In-house growth (following Hadoop story)

Access, access, access...

All the data sit in the DFS

Trivial to use massive compute power

→ lots of new applications



But... everything has to be MR

Cast any computation as map-only job

MPI, graph processing, streaming, *launching web-servers!?!*

Popularization

Everybody wants Big Data

Insight from raw data is cool

Outside MS and Google, Big-Data == Hadoop

Hadoop as catch-all big-data solution (and cluster manager)



Not just massive in-house clusters

New deployment environments

Small clusters (10s of machines)

Public Cloud

New challenges?

New deployment challenges

Small clusters

Efficiency matter more than scalability

Admin/tuning done by mere mortals

Cloud




Untrusted users (security)

Users are paying (availability, reliability)


Users are unrelated to each other (performance isolation)






Classic MapReduce



Large Shared Cluster



Classic MapReduce



Problems to solve

Who runs?

Where?


How much resources?

Order of execution

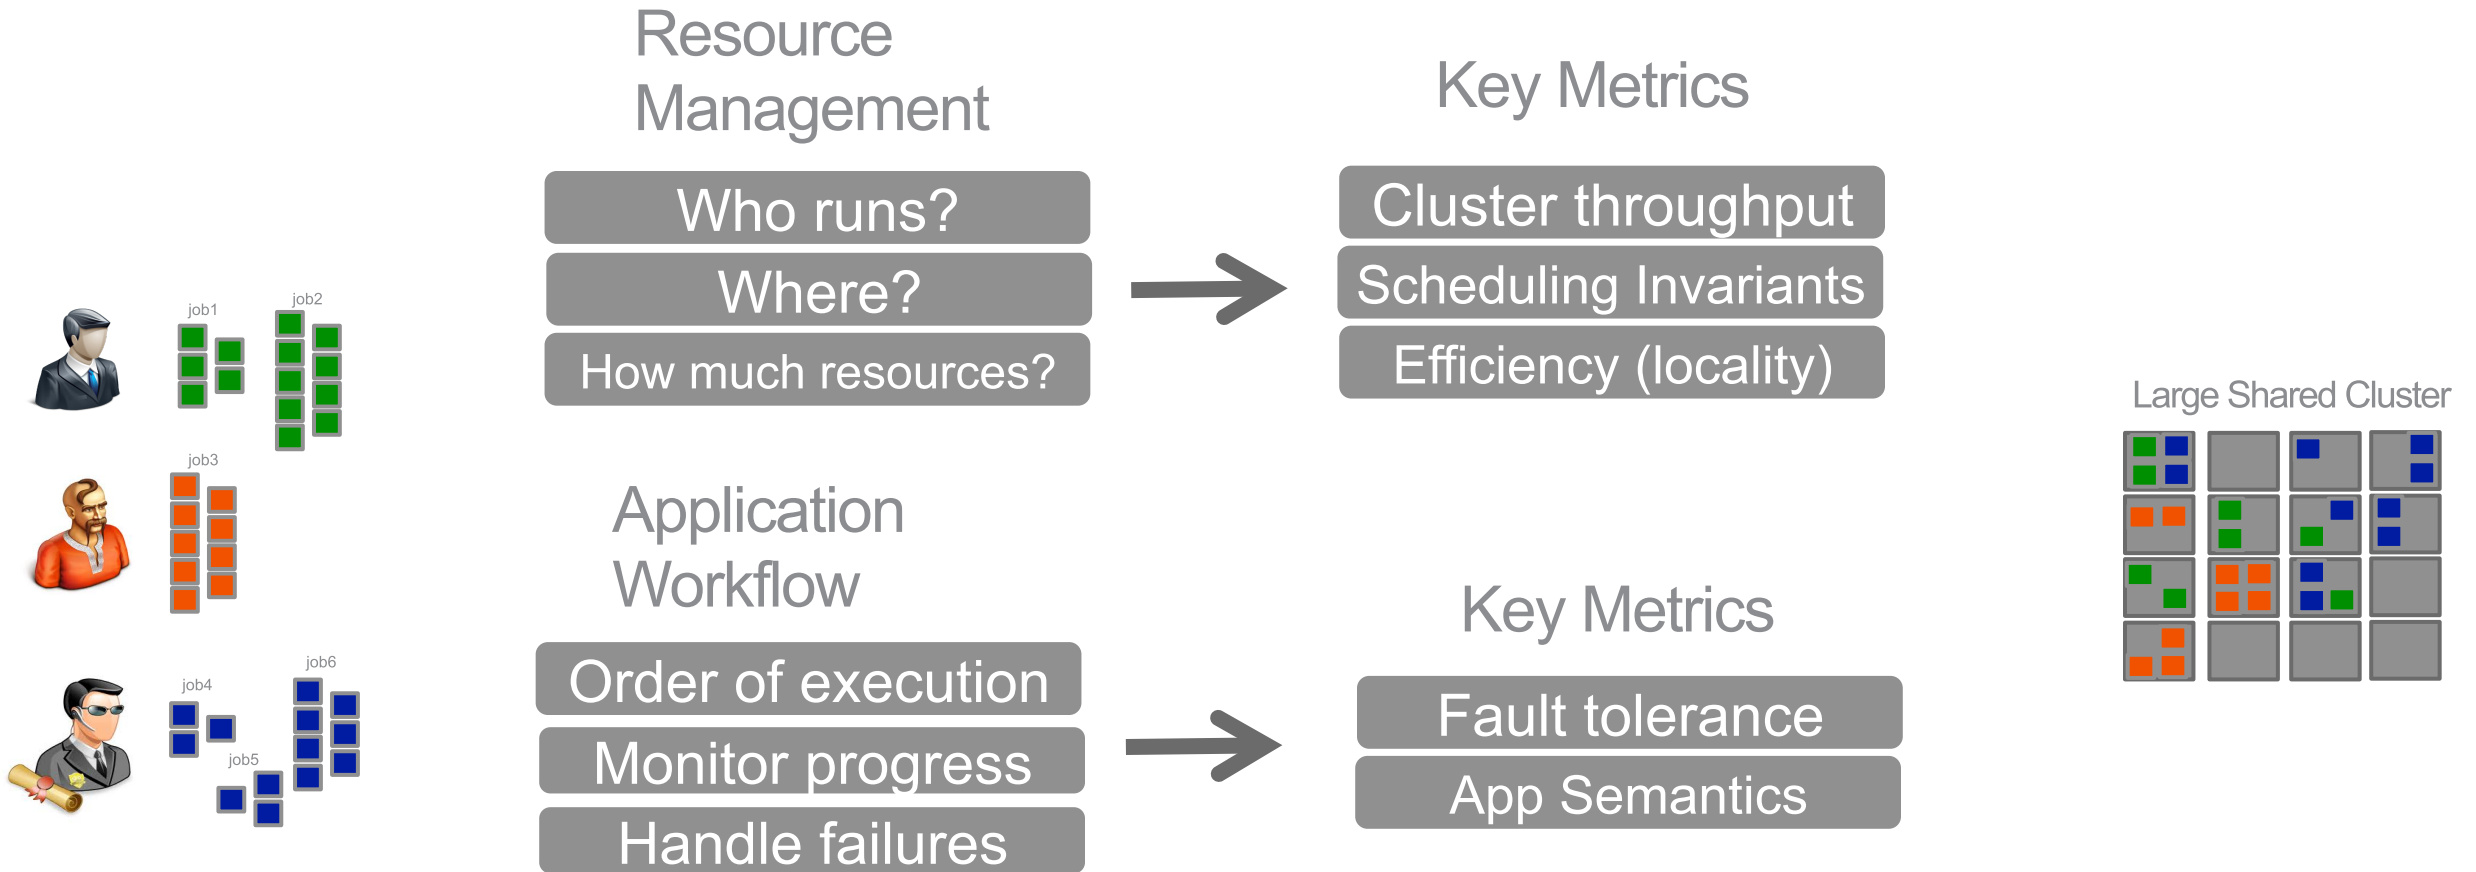
Monitor progress

Handle failures

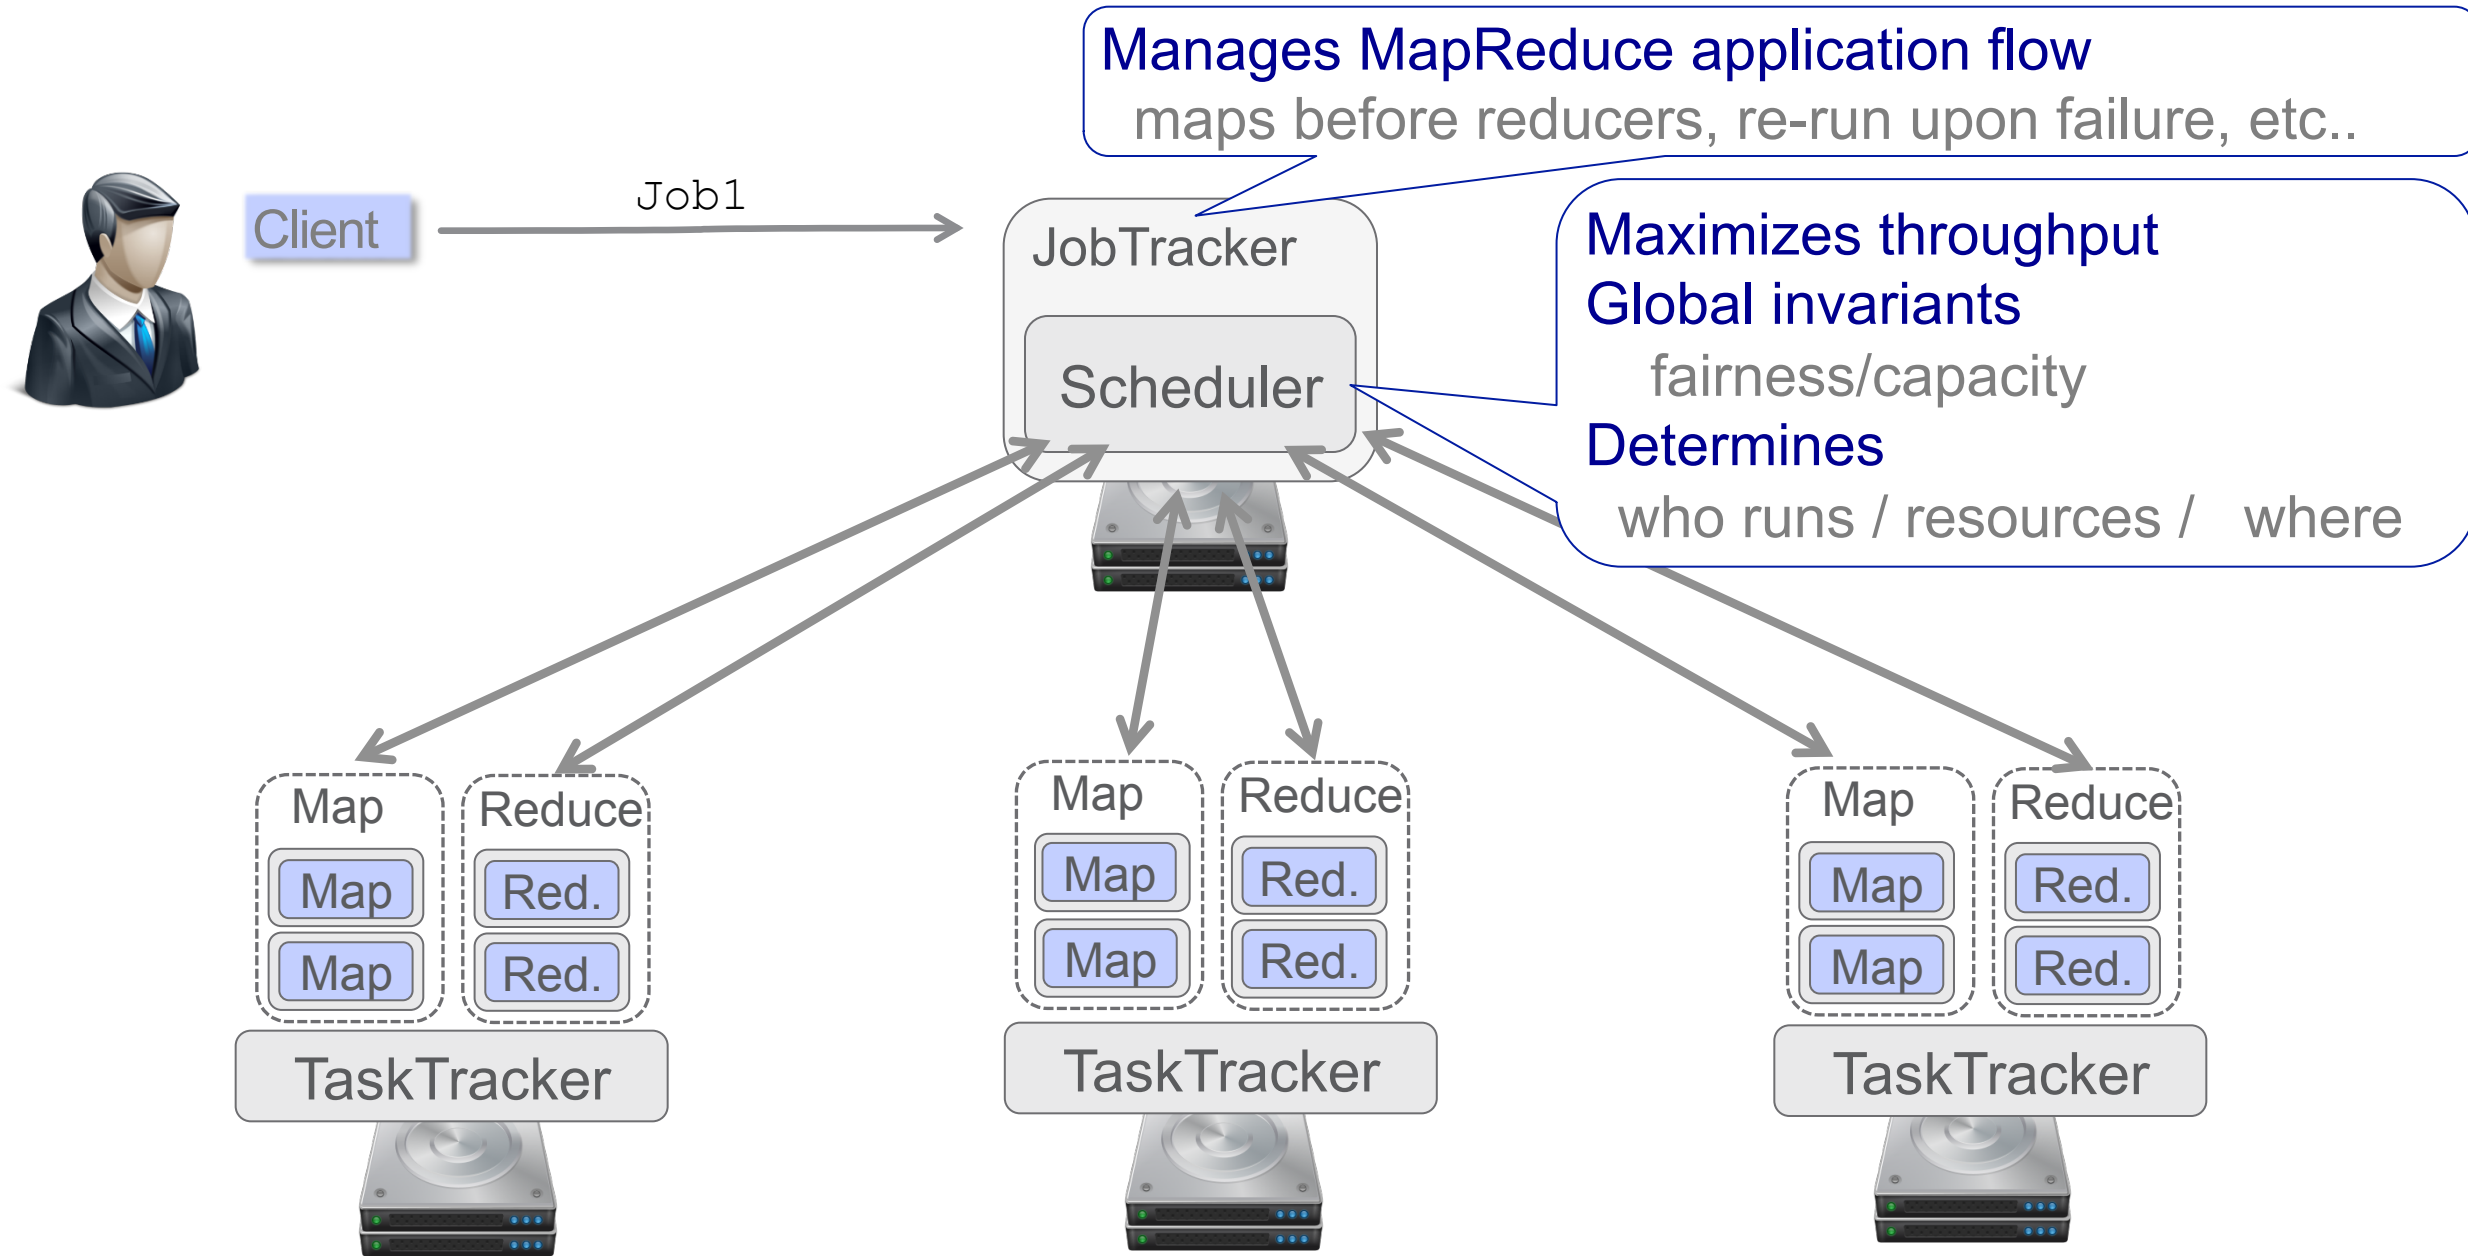
Large Shared Cluster



Classic MapReduce



Classic Hadoop (1.0) Architecture



Hadoop 1.0 Shortcomings

What are the key shortcomings
of (old) Hadoop?

Hadoop 1.0 Shortcomings (similar to original MR)

Programming model rigidity

JobTracker manages resources

JobTracker manages application workflow (data dependencies)

Performance and Availability

Map vs Reduce slots lead to *low cluster utilization* (~70%)

JobTracker had too much to do: *scalability concern*

JobTracker is a single point of failure

Cluster OS

“...towards general purpose computing...”

Three proposals

YARN (2008-2013, Hadoop 2.x, production at Yahoo!, GA)

Request-based central scheduler

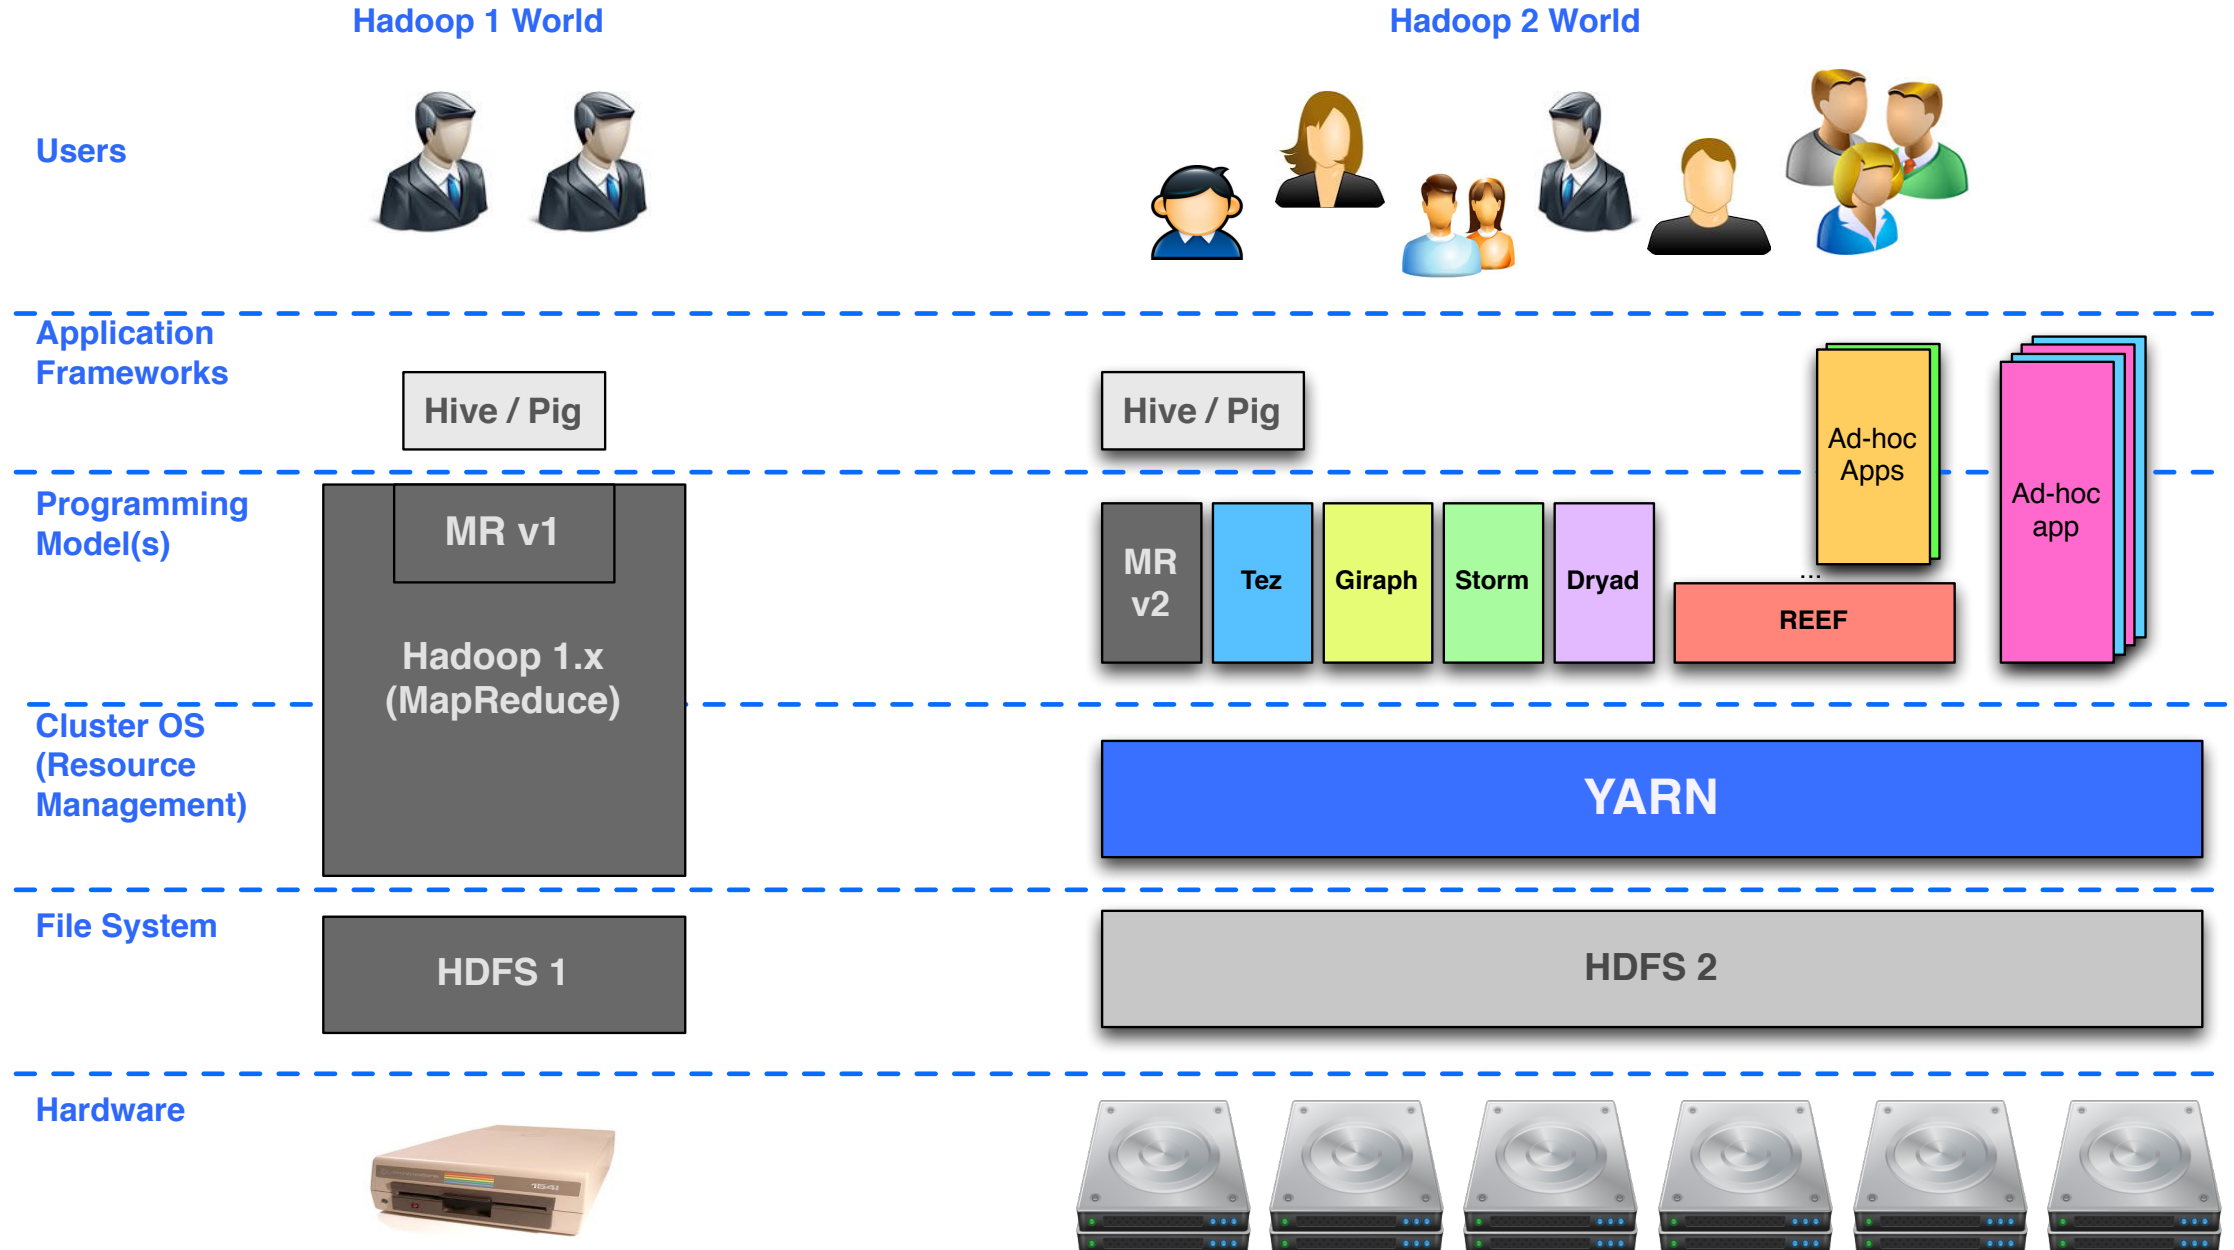
Mesos (2011, UCB, open-sourced, tested at Twitter)

Offer-based two level scheduler

Omega (2013, Google, simulation)

Shared-state-based scheduling

YARN

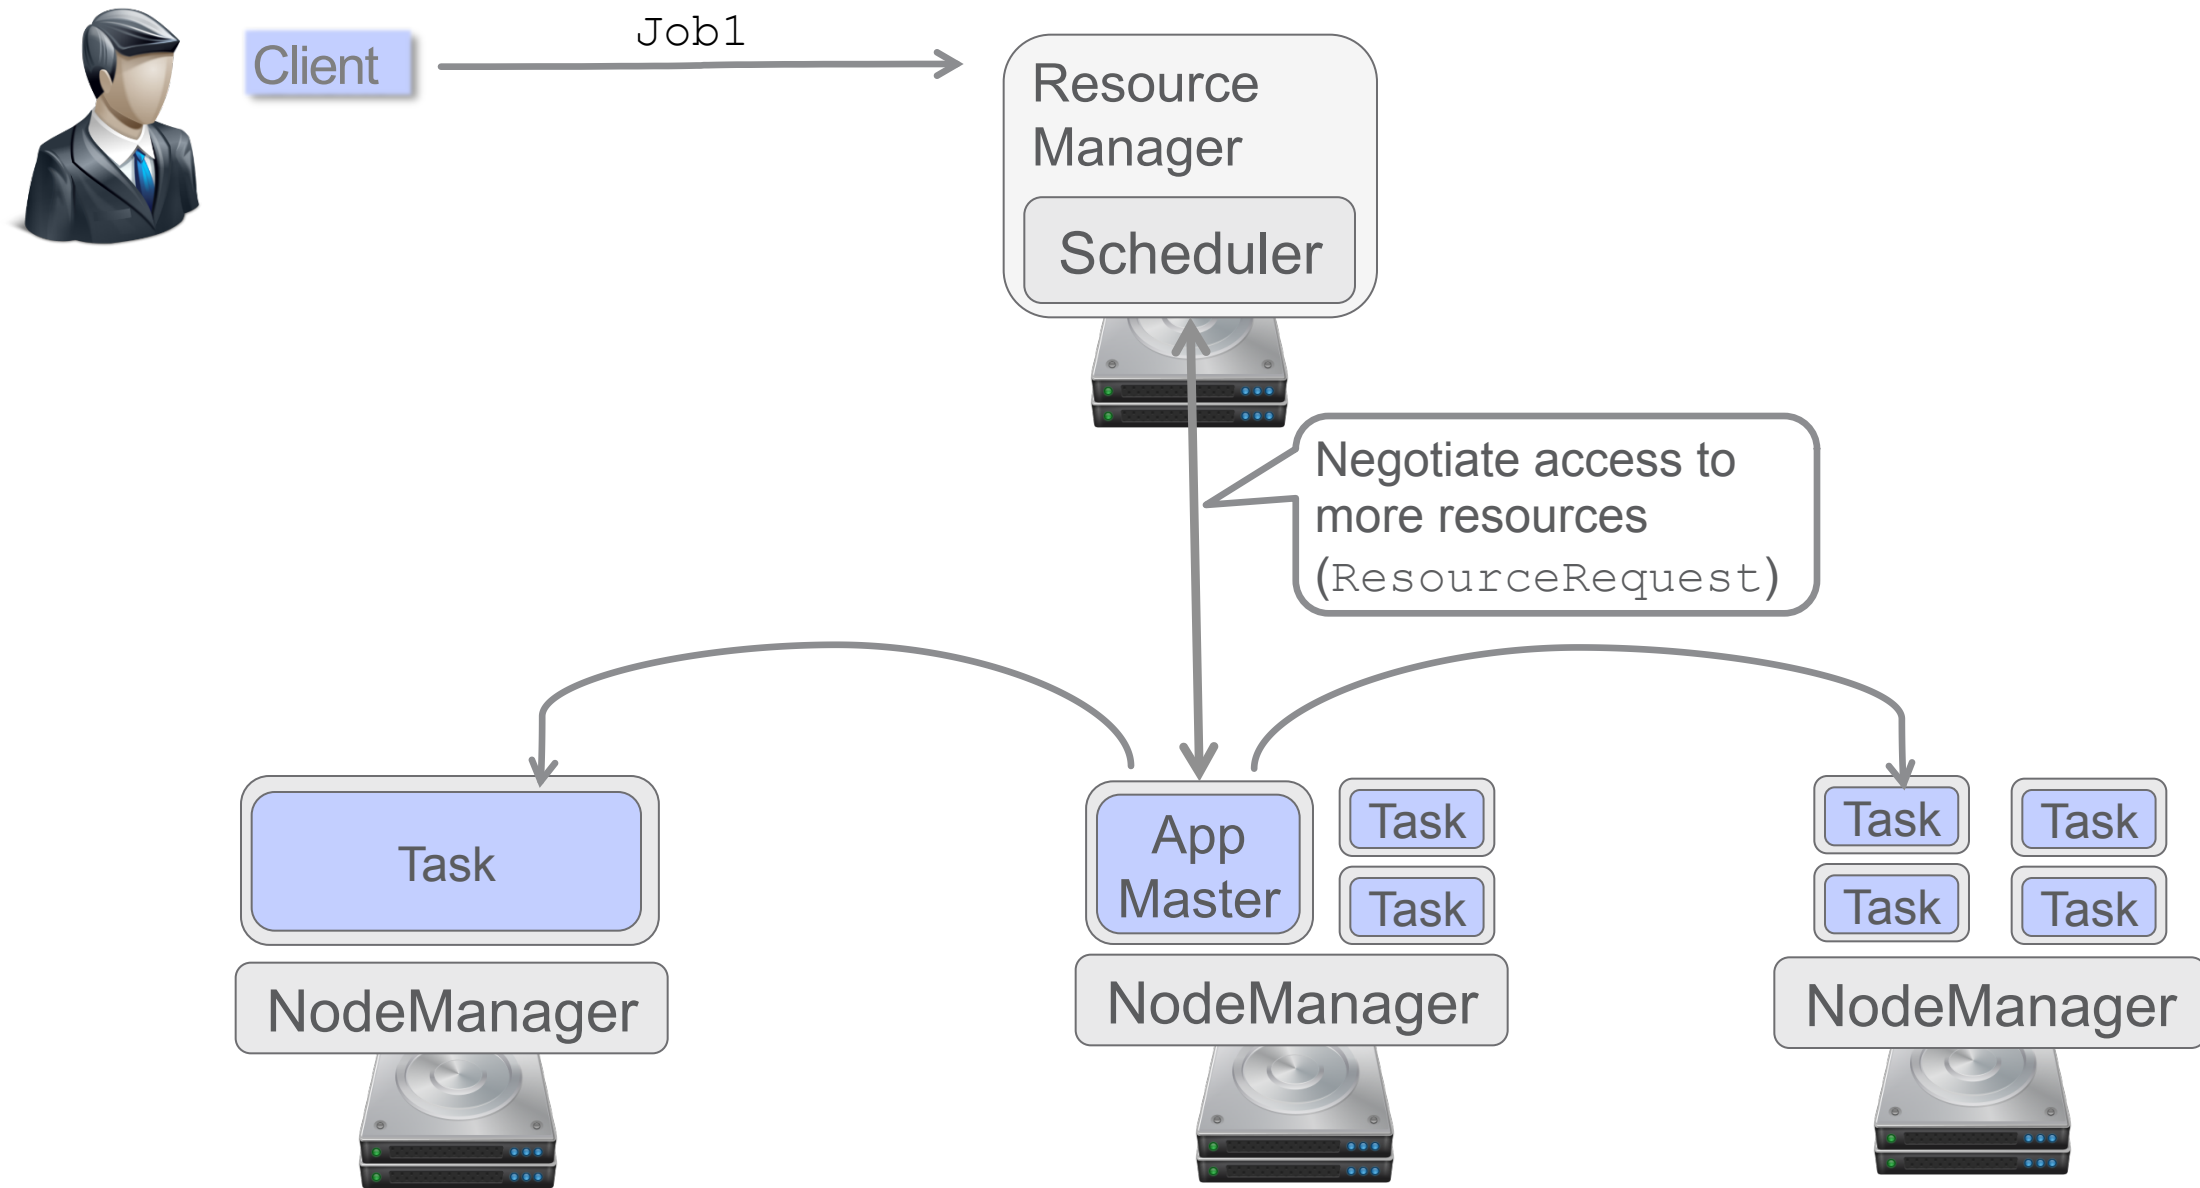


YARN (or Hadoop 2.x)

A new architecture for Hadoop


Decouples resource management from programming model
(MapReduce is an “application” running on YARN)

YARN (Hadoop 2) Architecture

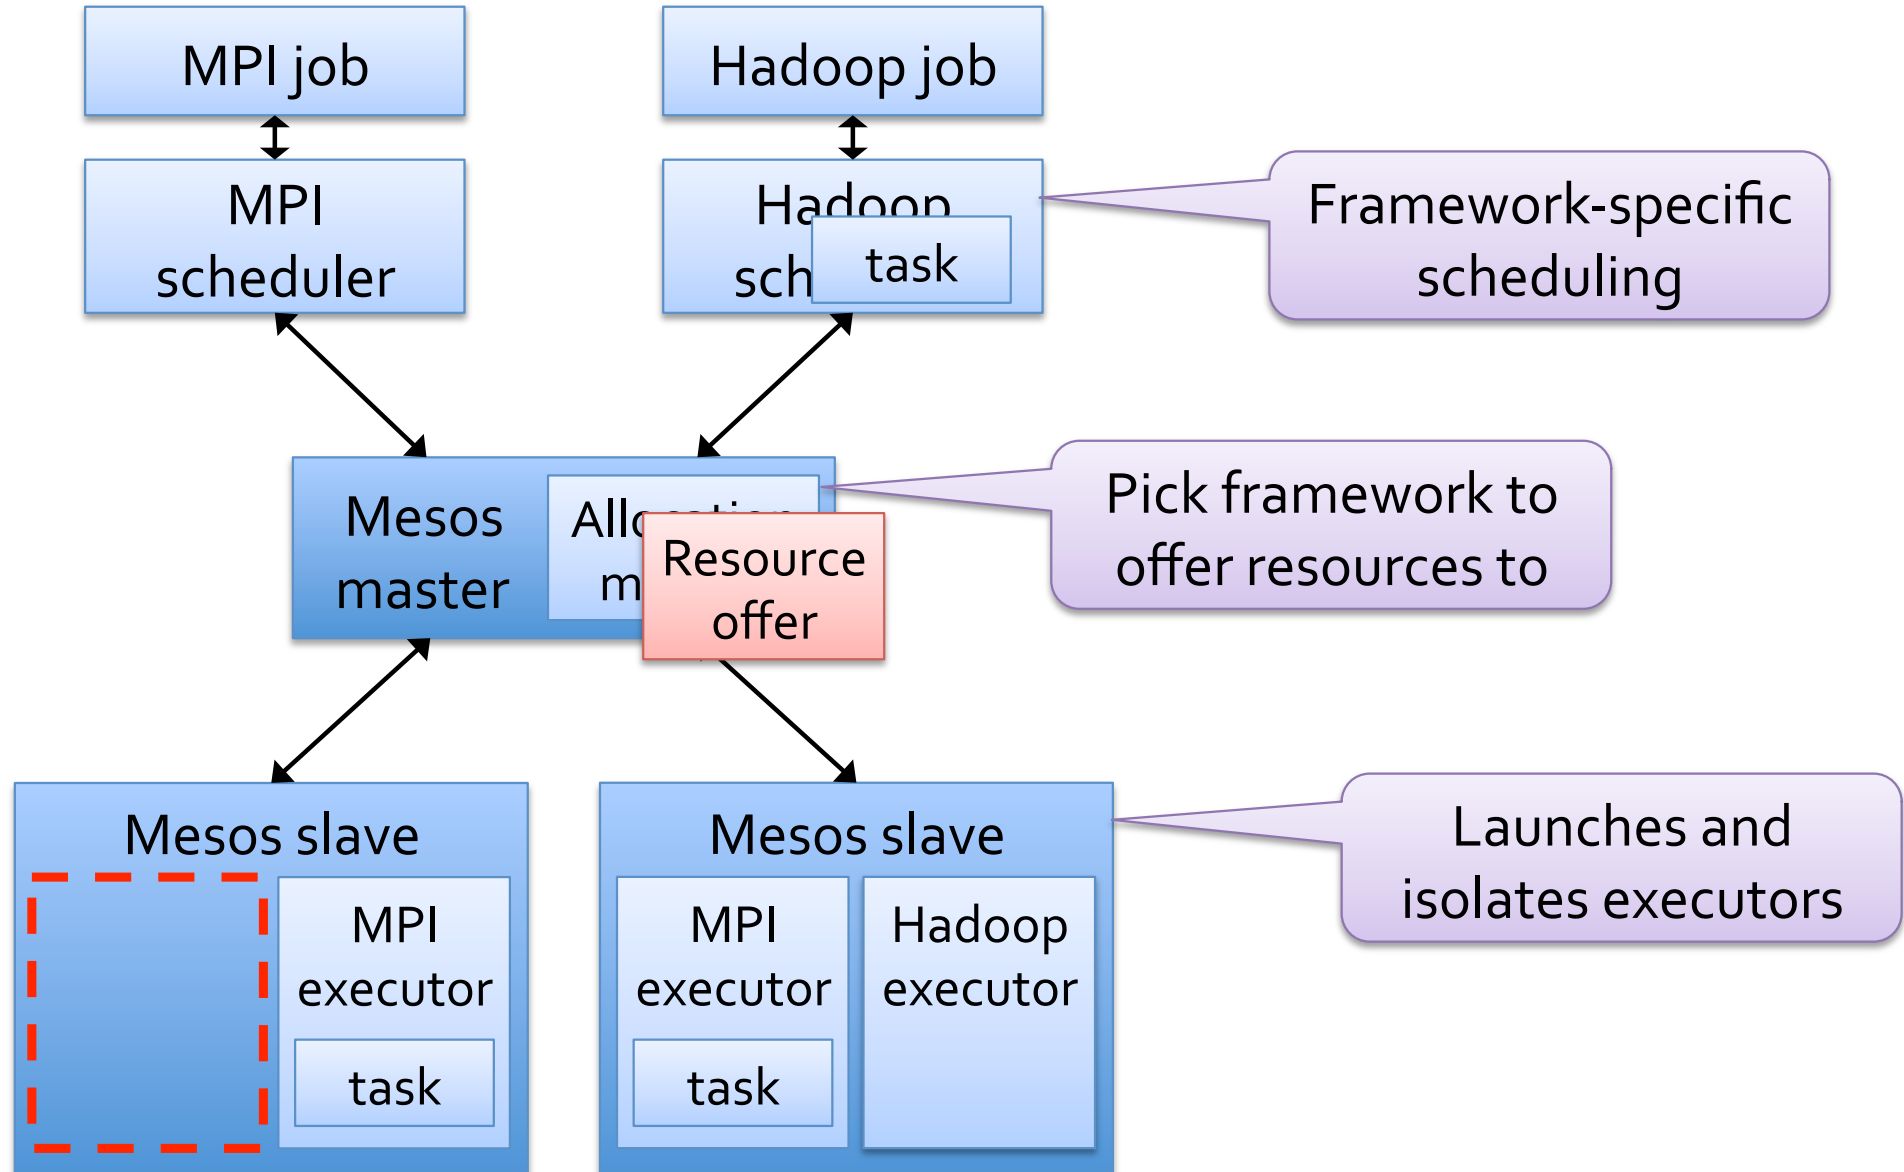


Mesos

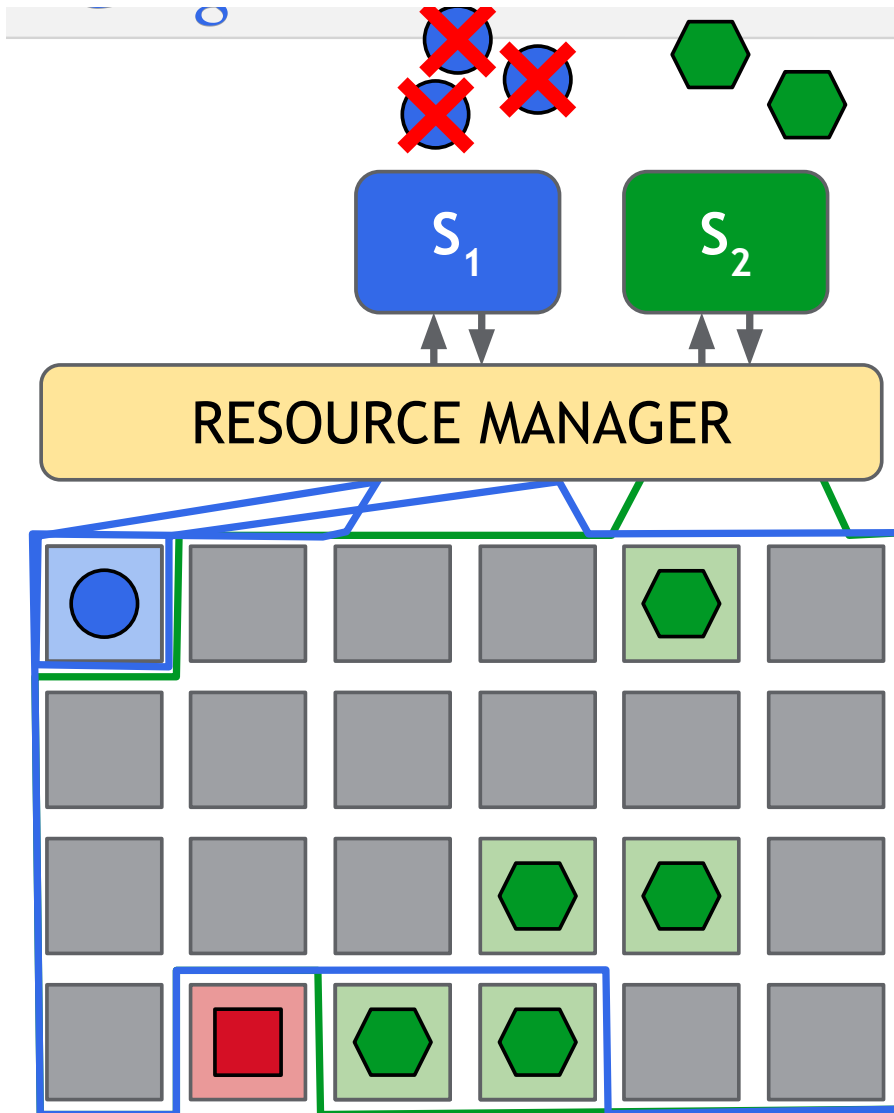
Layer between framework and hardware



Mesos



Omega critique of Mesos



1. Green receives offer of all available resources.

2. Blue's task finishes.

3. Blue receives tiny offer.

4. Blue cannot use it.

[repeat many times]

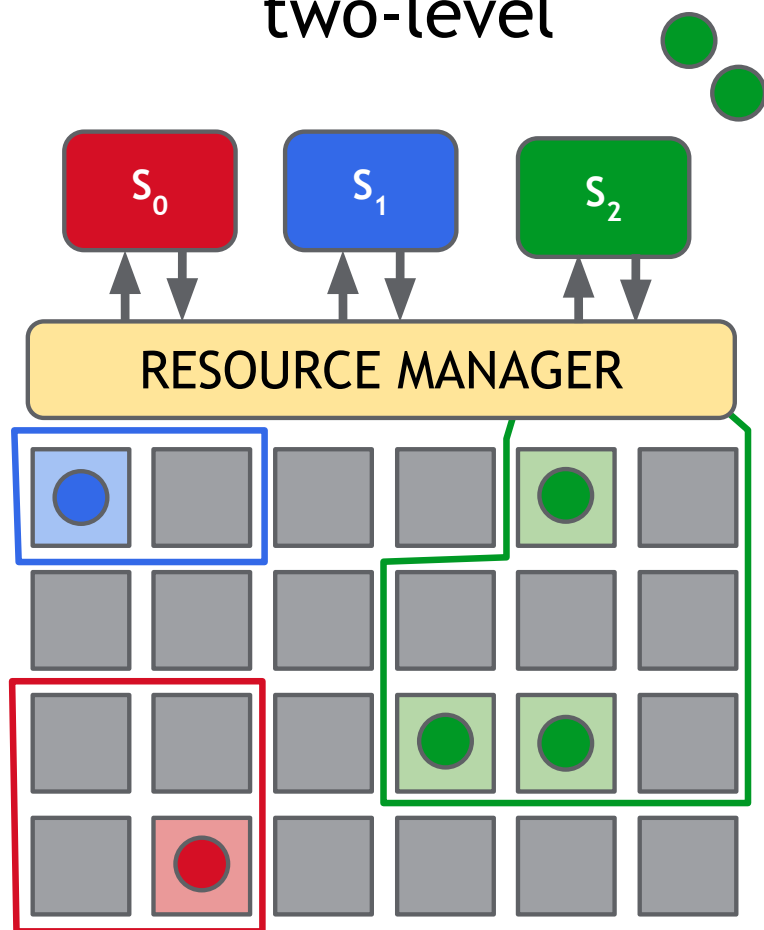
5. Green finishes scheduling.

6. Blue receives large offer.

By now, it has given up.


Omega

two-level




- hoarding
- information hiding

shared-state



Omega



Why does this matter?

Flexibility, Performance and Availability

Multiple Programming Models

Central components do less → *scale better*

Easier High-Availability (e.g., RM vs AM)

Anything else you can think?

Anything else you can think?

Maintenance, Upgrade, and Experimentation

Run with multiple framework versions (at one time)

Trying out a new ideas is as is as launching a job

Summarizing (technical considerations)

- ALL** improve scalability, flexibility, availability
- YARN** focus on “jobs”, untrusted apps, no faith sharing, global invariants are king, does not show cluster state
- Mesos** good upgrade, high-availability story, works for services, expose cluster state
- Omega** even more focus on scale, major trust in frameworks, global invariants afterthought, services are key, expose cluster state

Summarizing (non technical considerations)

YARN open-source, GA (since yesterday), used in-production

Mesos open-source, tested in real companies

Omega prototype, Google-internal